

C2_00 Gate B Protected Review - Astra RC4

Clinic Scheduling - open only after the independent Gate B attempt is saved

Purpose: diagnose reasoning and repair a failed competency. This is not an answer sheet to copy into the Gate.

1. Objective fixture controls

These are invariant controls from the supplied workbook. They help you detect population/grain errors; they do not choose the business decision for you.

Clinic	Appointments	No-show	Status pattern
Central Clinic	26	100.0%	26 NS / 0 C / 0 done
East Clinic	26	0.0%	0 NS / 26 C / 0 done
North Clinic	26	0.0%	0 NS / 0 C / 26 done

2. Expected reasoning path

Decision contract

The decision has two parts: identify where missed appointments occur and decide whether reminder policy should change. Those are not the same claim.

Fixture diagnosis

After controlling one duplicate, 78 unique appointments remain. Central Clinic contains 26 no-shows; East Clinic contains 26 cancellations; North Clinic contains 26 completions, including eight rescheduled flags.

Denominator logic

For a no-show rate, all appointments and "completed + no-show" are different eligible populations. State whether cancellations belong in the denominator. East has no completed/no-show rows, so a completed+no-show denominator is undefined there rather than zero.

Causal boundary

The extract does not prove reminder timing caused Central no-shows. Patient communications explicitly needs evidence connecting reminder timing to the issue. Recommend a reminder-policy test only as a next experiment unless additional evidence exists.

Validation

Reconcile clinic counts to 78 unique appointments; inspect rescheduled-ID semantics; compare no-show volume and rate; test whether the ranking changes under approved denominator choices.

3. Critical failure patterns

- Calling East "best" on no-show rate while ignoring that its records are all cancellations.
- Dividing by a zero/undefined eligible population and reporting 0%.
- Treating rescheduled appointments as independent without resolving ID semantics.
- Claiming reminder timing caused no-shows from this extract.
- Changing policy without a decision threshold or test plan.

4. Competency routing and repair proof

Competency	Pass evidence
Decision/scope	Separate diagnosis from policy decision; specify threshold and audience.
Metric contract	Define no-show numerator and eligible population; handle cancellations/reschedules.
Grain/source	One row = appointment event/ID under stated system behavior; duplicate controlled.
Judgment	Association vs causal policy claim separated.
Validation	Clinic/overall counts reconcile; undefined denominators handled explicitly.
Handoff/defense	Recommendation includes test/owner/timing and evidence that would reverse it.

PASS PHILOSOPHY: No percentage gaming. One unresolved material definition, grain, duplicate, validation, or evidence-strength failure blocks the Gate even if other sections are polished.

5. Changed repair proof

Changed retry after repair: a new outpatient dataset mixes completed, late-cancelled and no-show visits. Define an eligible population, identify the operational problem, and propose a reminder experiment without claiming causality.

NEXT ROUTE: Gate C only after targeted repair and explicit assignment.