

COURSE 2A

C2-07 CLOUD / FABRIC / ENTERPRISE BI

BEGINNER-FIRST TEACHING BOOK - RC4

Cloud boundaries -> Fabric -> Service -> deployment -> refresh -> governance -> monitoring

THIS BOOK IS YOUR TEACHER

Platform-specific references refreshed 2026-09-10. Internal concepts remain sufficient if a link or trial is unavailable.

This book is your teacher

Use one lesson at a time. New advanced concepts are still taught from first principles. Save a genuine attempt before opening protected review; do not skim ahead just because a topic name looks familiar.

Daily lesson loop

1) Why it matters -> 2) what you learn -> 3) prerequisite -> 4) picture it in your head -> 5) new words -> 6) watch / visual source when useful -> 7) follow with me once -> 8) expected result -> 9) why it worked -> 10) common beginner mistakes -> 11) recovery ladder -> 12) small fresh task -> 13) validate -> 14) teach it back -> 15) protected review + changed retry -> 16) evidence / mastery + exact next route.

Source rule

External videos are companions, not the course authority. If a useful source exists, open the exact link and return to the course task. If an external source would leak a solved capstone or add little value, the lesson explicitly uses an INTERNAL TEACHING ROUTE.

Attempt protection

Do the teacher example once. Then close it and complete the changed task from your own blank workbook/query/report/notebook. Protected review opens only after a saved attempt and is followed by a changed retry.

Mastery standard

You are finished only when you can reproduce the result, validate it, explain why it works, recover from a mistake and save the required evidence.

Contents / Index

Every canonical C2-07 lesson and actual page number

This book is your teacher	2
Contents / Index	3
Prerequisite and tool map	4
Practice, review and Gate route	5
CLD1 Cloud data fundamentals	6
CLD2 Compute, storage and identity concepts	10
CLD3 Microsoft Fabric ecosystem	14
CLD4 OneLake, lakehouse and warehouse concepts	18
CLD5 Power BI Service in enterprise environments	22
CLD6 Workspaces and deployment concepts	26
CLD7 Refresh, gateways and scheduled operations	30
CLD8 Permissions and governance	34
CLD9 Monitoring and failure awareness	38
CLD10 Analyst - analytics engineer - data engineer collaboration	42
Module mastery checklist + quick reference	46
Source / video registry	47

Prerequisite and tool map

What is assumed, what is new and what can be environment-blocked

ASSUMED

Course 1 Power BI basics; C2-04 lifecycle/RLS/reporting; C2-06 warehouse/lakehouse/lineage concepts. New cloud/Fabric/enterprise-operating terms are defined before use.

CORE TOOLS

Teaching Book + C2-07 Enterprise Operations Workbook + synthetic item/workspace/source/job/access fixtures. These are enough for required reasoning and Gate evidence.

OPTIONAL LIVE TOOLS

Power BI Service / Microsoft Fabric tenant or trial, deployment pipelines, gateway, Monitor and OneLake catalog. Availability depends on tenant, license, capacity, region and organizational policy.

PL-300 POSITION

CLD5-CLD8 directly reinforce current Manage and Secure Power BI skills: workspaces/apps/distribution, gateways/refresh, roles/item/model access/RLS and sensitivity labels. Course 2A remains optional; PL-300 is supplementary, not a hidden DA-to-DE prerequisite.

Practice, review and Gate route

Architecture plan -> operational evidence -> incident/change transfer



LESSON RULE

Use matching P_CLD# workbook sheet after the exact lesson pages. Save first enterprise plan/evidence before protected review. Close the review before changed transfer.

MINI-LAB

NorthStar Enterprise Analytics Deployment: workspaces/items, builder/consumer access, DEV/TEST/PROD, refresh/gateway, governance, monitoring incident and team handoff.

GATES

A Fleet Maintenance, B Education Enrollment, C Energy Metering. Distinct workspace/source/access/refresh/incident fixtures. Any critical unsafe access or unowned failure blocks pass.

CLD1 - Cloud data fundamentals

1. Why it matters

Cloud analytics can feel like a wall of product names. Before Fabric, you need a simple model: who owns the service, where your organization lives, where teams work, what provides compute and where analytical items/data are organized.

2. What you are learning

Explain cloud responsibility at a beginner level; distinguish tenant, workspace, capacity and item; recognize Fabric as SaaS and avoid confusing 'cloud' with 'someone else's laptop'.

3. Prerequisite / what you already know

C2-06 workload/warehouse concepts help, but no Azure administration experience is assumed.

4. Picture it in your head

Mental model

Apartment utility model: the organization/tenant owns the environment and identities; workspaces organize team content; capacity supplies managed compute; items/data are the things teams build and govern. Cloud removes infrastructure chores, not ownership, security, cost or validation responsibility.

5. New words before we use them

SaaS: Software delivered and operated as a managed service; you use/configure it rather than managing all infrastructure. Tenant: Organization-level boundary for identities, settings and Fabric/Power BI environment. Capacity: Allocated compute resources used by Fabric workloads; separate concept from OneLake storage.

6. Watch / visual source

VISUAL/SOURCE ROUTE: Microsoft Fabric overview is required visual/reference. Trial is optional and never required for mastery.

<https://learn.microsoft.com/en-us/fabric/get-started/microsoft-fabric-overview>

After watching/opening: return to this lesson and reproduce the worked example from your own file/query/report. Save course-task evidence, not a screenshot of the source.

7. Follow with me once

Teacher demonstration - do this once with me, then close the example before your fresh task.

- 1 Business asks for governed Sales analytics used by 200 employees. Start with the organization/tenant-not with a VM.
- 2 Create a conceptual Analytics workspace for team-owned Fabric/Power BI items. A workspace is a collaboration/security boundary inside the tenant.
- 3 Place a semantic model/report and, if needed, a lakehouse/warehouse/pipeline inside the workspace.
- 4 Capacity supplies compute for Fabric workloads; identity and tenant policies control who can use/create items. OneLake is Fabric's unified logical data lake, not the same thing as capacity.
- 5 Write which responsibilities remain with Microsoft (managed service/infrastructure) versus your organization (access, data, workspace design, definitions, governance, cost/capacity decisions).

8. What you should see / be able to say

Expected result

You can place tenant -> workspace -> items and explain capacity/identity responsibilities without Azure-infrastructure jargon.

9. Why it worked

Reasoning

The mental model separates organization, collaboration boundary, compute and data/items instead of treating 'cloud' as one box. Professional upgrade: Fabric is currently delivered as SaaS, and Fabric tenants automatically include OneLake. Capacity sizing/licensing can change, so this course teaches responsibilities and boundaries before SKU memorization.

10. Common beginner mistakes

1. Tenant = workspace. | 2. Capacity = storage folder. | 3. Cloud removes responsibility for permissions/governance. | 4. Assuming Fabric requires you to provision Azure resource groups for ordinary use.

11. If you get stuck - recovery ladder

Recovery

Draw four labeled boxes only: Organization/Tenant -> Workspace -> Items, with Capacity beside them as compute. Add OneLake as shared storage foundation. Save the genuine attempt before review.

12. Now you do a small version alone

Fresh independent task

Fresh case: a retail company needs Finance and Operations analytics with separate ownership but some shared certified data. Sketch tenant/workspace/capacity/item boundaries.

13. Validate your result

Validation

Your diagram names owners, access boundary, compute boundary and shared-data path; no product term is used without a job.

14. Teach it back in your own words

Explain without reading

Explain tenant, workspace and capacity to a manager using organization -> team room -> engine.

15. Protected review + changed fresh retry

Attempt protection

Save first diagram. Changed retry uses a consulting company with two client workspaces and a central governed semantic-model team.

16. Evidence / mastery + exact next route

Mastery evidence

Cloud responsibility matrix + tenant/workspace/capacity/item diagram + explain-back. After saving evidence, continue to CLD2.

CLD2 - Compute, storage and identity concepts

1. Why it matters

Most enterprise incidents happen because someone confuses 'can the engine run?' with 'is this user allowed?' or 'where does the data live?'. Compute, storage and identity are separate control planes.

2. What you are learning

Distinguish compute, storage and identity; apply least privilege to people/groups/service principals conceptually; explain why a job may have compute but still fail due to data/credential permission.

3. Prerequisite / what you already know

CLD1 tenant/workspace/capacity model.

4. Picture it in your head

Mental model

Kitchen model: storage is the pantry that keeps ingredients, compute is the stove that performs work, and identity is the keycard that decides who may enter and use each resource. Data can persist when compute stops, and available compute does not override access policy.

5. New words before we use them

Compute: Resources that execute queries, refreshes, pipelines or notebooks. Identity: User, group or application/service principal that requests access. Least privilege: Grant only the minimum permissions needed for the intended job, for the required scope/time.

6. Watch / visual source

VISUAL/SOURCE ROUTE: Book-led; current workspace-role and Fabric overview references are enough.

<https://learn.microsoft.com/en-us/fabric/fundamentals/roles-workspaces>

After watching/opening: return to this lesson and reproduce the worked example from your own file/query/report. Save course-task evidence, not a screenshot of the source.

7. Follow with me once

Teacher demonstration - do this once with me, then close the example before your fresh task.

- 1 Capacity is available: the service has compute to run a semantic-model refresh.
- 2 Refresh identity/connection still needs permission and valid credentials to reach its data source.
- 3 The workspace role controls what a person/service principal can do with Fabric/Power BI items, but source-system permission can be separate.
- 4 OneLake/item permissions and semantic-model/RLS permissions can further constrain data access.
- 5 Create a matrix: Identity -> Resource -> Needed action -> Minimum permission -> Owner. Do not solve access by granting Admin everywhere.

8. What you should see / be able to say

Expected result

A clear diagnosis separating compute availability, item permission, source credentials and data-level security.

9. Why it worked

Reasoning

Each failure is tested at the control plane where it can actually occur. Professional upgrade: Enterprise automation increasingly uses groups and service principals rather than personal accounts. Exact tenant/API policies vary; the durable practice is explicit identity/action/scope/owner.

10. Common beginner mistakes

1. Admin workspace role fixes every source-access problem. | 2. Healthy capacity proves data source is reachable. | 3. RLS replaces all workspace/item permissions. | 4. Use personal credentials for production because they are easiest.

11. If you get stuck - recovery ladder

Recovery

For any 'access failed' issue ask four questions: Who is the identity? Which resource? Which action? Which permission/credential controls that action? Save the genuine attempt before review.

12. Now you do a small version alone

Fresh independent task

Fresh case: service principal publishes items, analyst builds reports, viewer consumes app, gateway identity reaches SQL Server. Design minimum conceptual permissions.

13. Validate your result

Validation

No identity receives unexplained broad Admin access; source, workspace and data-level permissions are separate.

14. Teach it back in your own words

Explain without reading

Explain why capacity, permission and data credentials are three different reasons a cloud job can succeed or fail.

15. Protected review + changed fresh retry

Attempt protection

Attempt first. Changed retry removes one employee and requires group-based access without editing every item manually.

16. Evidence / mastery + exact next route

Mastery evidence

Identity-action-permission matrix + failure diagnosis + explain-back. After saving evidence, continue to CLD3.

CLD3 - Microsoft Fabric ecosystem

1. Why it matters

Fabric combines many analytics experiences. Memorizing icons does not help you design a solution; you need to know which workload solves ingestion, engineering, warehouse, real-time or BI consumption problems and how they share OneLake.

2. What you are learning

Map major Fabric workloads to responsibilities; describe Fabric as end-to-end SaaS analytics; choose the smallest set of workloads needed for a business flow.

3. Prerequisite / what you already know

CLD1-2 and C2-06 architecture basics.

4. Picture it in your head

Mental model

Shared analytics campus: Fabric workloads are specialized rooms around shared OneLake and common governance. Data Factory moves/orchestrates, engineering transforms, warehouse serves SQL analytics, real-time handles events and Power BI serves semantic/report consumption; use only the rooms the solution needs.

5. New words before we use them

Workload: Fabric experience optimized for a class of analytics tasks. Data Factory: Fabric workload/capabilities for data movement and orchestration. Real-Time Intelligence: Fabric workload for ingesting/analyzing event and streaming data.

6. Watch / visual source

VISUAL/SOURCE ROUTE: Required: current What is Microsoft Fabric overview; optional Power BI Meets Fabric video.

<https://learn.microsoft.com/en-us/fabric/get-started/microsoft-fabric-overview>

After watching/opening: return to this lesson and reproduce the worked example from your own file/query/report. Save course-task evidence, not a screenshot of the source.

7. Follow with me once

Teacher demonstration - do this once with me, then close the example before your fresh task.

- 1 Data arrives from CRM/order sources. Data Factory may orchestrate copy/ingestion.
- 2 Data Engineering may use Spark/lakehouse transformations when that tool fits the workload.
- 3 Data Warehouse may serve governed relational dimensional tables queried with T-SQL.
- 4 Real-Time Intelligence is appropriate when operational event streams require real-time analysis; do not add it to a daily batch just because it exists.
- 5 Power BI provides semantic/reporting consumption. All required workloads can coexist in Fabric and access shared OneLake data/items according to permissions.

8. What you should see / be able to say

Expected result

A workload map where every Fabric component has a reason and unnecessary services are omitted.

9. Why it worked

Reasoning

You design from task -> workload instead of icon -> task. Professional upgrade: Microsoft currently lists Data Engineering, Data Factory, Data Science, Real-Time Intelligence, Data Warehouse, Databases and Power BI among Fabric experiences. The list can evolve; workload-to-job reasoning is the durable skill.

10. Common beginner mistakes

1. Every solution needs every Fabric workload. | 2. Data Factory = the warehouse. | 3. Power BI Service and Fabric are unrelated platforms. | 4. Real-Time Intelligence for monthly CSV reporting.

11. If you get stuck - recovery ladder

Recovery

Write the verbs: ingest, transform, store/query, stream, model/report. Map one tool only when a verb needs it. Save the genuine attempt before review.

12. Now you do a small version alone

Fresh independent task

Fresh case: IoT maintenance events + daily ERP finance + executive BI. Choose Fabric workloads for each path and explain integration.

13. Validate your result

Validation

Every chosen workload solves a named problem; no duplicate/unnecessary workload is added.

14. Teach it back in your own words

Explain without reading

Explain Fabric to a beginner as one managed analytics platform containing role-specific work areas over shared data.

15. Protected review + changed fresh retry

Attempt protection

Attempt first. Changed retry removes streaming need and simplifies the design.

16. Evidence / mastery + exact next route

Mastery evidence

Workload decision map + simplified alternative + explain-back. After saving evidence, continue to CLD4.

CLD4 - OneLake, lakehouse and warehouse concepts

1. Why it matters

C2-06 taught when a warehouse/lakehouse pattern fits. Here the new question is how Fabric organizes and shares those items over OneLake and when shortcuts/shared storage can reduce unnecessary data copies.

2. What you are learning

Explain OneLake tenant-wide logical storage, workspaces/items and shortcuts; place lakehouse/warehouse items in an enterprise Fabric topology; avoid duplicating C2-06 dimensional-design theory.

3. Prerequisite / what you already know

C2-06 DWF8-DWF9 plus CLD1-3.

4. Picture it in your head

Mental model

One storage estate with multiple serving patterns: OneLake is the tenant-wide logical lake; workspaces organize ownership; shortcuts can reference supported data without traditional copying; lakehouse favors file/table engineering patterns while warehouse favors governed relational SQL serving.

5. New words before we use them

OneLake: Fabric's unified logical data lake for the organization. Shortcut: Reference that provides access to data in another supported location without traditional copy/migration. Item: Fabric object such as lakehouse, warehouse, semantic model or pipeline inside a workspace.

6. Watch / visual source

VISUAL/SOURCE ROUTE: Required visual: OneLake overview; optional 'No Walls, Just Data' session.

<https://learn.microsoft.com/en-us/fabric/onelake/onelake-overview>

After watching/opening: return to this lesson and reproduce the worked example from your own file/query/report. Save course-task evidence, not a screenshot of the source.

7. Follow with me once

Teacher demonstration - do this once with me, then close the example before your fresh task.

- 1 Central Data Platform workspace owns curated customer/sales data in a Warehouse or Lakehouse depending on workload.
- 2 Finance BI workspace consumes governed semantic/reporting assets rather than receiving uncontrolled CSV copies.
- 3 Where appropriate, OneLake shortcuts can expose existing data without duplicating it into another lakehouse; still apply ownership/security.
- 4 Document which workspace owns the authoritative item and which team only consumes it.
- 5 Do not use OneLake to erase governance boundaries: shared storage does not mean every user sees every item.

8. What you should see / be able to say

Expected result

A Fabric topology with clear data ownership, shared use and minimized unnecessary copies.

9. Why it worked

Reasoning

OneLake unifies storage access while workspaces/items/security preserve organizational control. Professional upgrade: OneLake currently supports unified tenant storage and zero-copy/shortcut patterns to supported external sources. Exact shortcut/security capabilities can evolve, so design must be validated against current docs before production.

10. Common beginner mistakes

1. OneLake = one giant folder everyone can read. | 2. Shortcut automatically grants access to every consumer. | 3. Duplicate every dataset per workspace 'for safety'. | 4. Re-teaching warehouse/lakehouse choice without Fabric topology.

11. If you get stuck - recovery ladder

Recovery

Start with authoritative owner and consumer. Ask whether data must be copied, referenced or surfaced as a semantic model. Save the genuine attempt before review.

12. Now you do a small version alone

Fresh independent task

Fresh case: Marketing needs governed Customer 360 from central lakehouse while Finance uses a warehouse and shared certified semantic model. Draw item/workspace/data paths.

13. Validate your result

Validation

For each copy/reference, state why it exists; ownership and permissions remain explicit.

14. Teach it back in your own words

Explain without reading

Explain OneLake as one organizational data-lake foundation, not one unrestricted spreadsheet folder.

15. Protected review + changed fresh retry

Attempt protection

Attempt first. Changed retry adds an external S3/ADLS source and asks whether a shortcut is preferable to copy.

16. Evidence / mastery + exact next route

Mastery evidence

Workspace/item topology + ownership/copy decision + explain-back. After saving evidence, continue to CLD5.

CLD5 - Power BI Service in enterprise environments

1. Why it matters

Publishing a PBIX is not the same as enterprise delivery. Teams need clear builder and consumer routes, trusted semantic models, distribution methods, ownership and support.

2. What you are learning

Design workspace collaboration, semantic-model/report ownership, app/direct share choices, content endorsement and consumer access; separate build surface from polished distribution.

3. Prerequisite / what you already know

C2-04 DAX15 lifecycle basics and CLD1-4.

4. Picture it in your head

Mental model

Publishing supply chain: builders work in governed workspaces, semantic models define reusable analytical meaning, reports present it, and apps/direct sharing distribute it to intended audiences. Production delivery also needs ownership, endorsement, refresh, security and support - not just a successful Publish click.

5. New words before we use them

App: Packaged Power BI distribution experience published from a workspace to audiences. Semantic model: Governed analytical model consumed by reports/Excel/other clients. Endorsement: Power BI/Fabric trust signal such as promoted/certified content under organizational policy.

6. Watch / visual source

VISUAL/SOURCE ROUTE: Optional PL-300 Manage and Secure video; current Microsoft collaboration docs control behavior.

<https://learn.microsoft.com/en-us/power-bi/collaborate-share/collaborate-share-overview>

After watching/opening: return to this lesson and reproduce the worked example from your own file/query/report. Save course-task evidence, not a screenshot of the source.

7. Follow with me once

Teacher demonstration - do this once with me, then close the example before your fresh task.

- 1 BI builders collaborate in a controlled workspace. Consumers do not automatically need edit access to the workspace.
- 2 Create/use a governed semantic model with documented owner/freshness/security and reports built on it.
- 3 Choose app for broad polished distribution; use direct sharing for narrower cases when appropriate. Record licensing/capacity/tenant constraints before real rollout.
- 4 Promote/certify trusted content according to organizational governance, rather than letting five teams publish competing 'official revenue' models.
- 5 Define support owner, refresh owner, change notice and consumer path.

8. What you should see / be able to say

Expected result

An enterprise delivery plan separating building, governed semantic truth, distribution and support.

9. Why it worked

Reasoning

Workspace collaboration permissions are not confused with consumer distribution. Professional upgrade: Current PL-300 directly measures creating/configuring workspaces/apps, publishing/updating items, distribution methods, subscriptions/alerts and content endorsement. Tenant/license constraints must be checked at implementation time.

10. Common beginner mistakes

1. Give all consumers Contributor because app setup is unfamiliar. | 2. Every report owns its own copied semantic model. | 3. Certified means mathematically correct forever. | 4. No support/change owner after publish.

11. If you get stuck - recovery ladder

Recovery

Draw Builders -> Workspace -> Semantic Model -> Report/App -> Consumers. Add owner/security/freshness at the model and support owner at distribution. Save the genuine attempt before review.

12. Now you do a small version alone

Fresh independent task

Fresh case: Sales has 6 builders, 40 managers and 600 viewers. Design workspace + semantic model + app/share route.

13. Validate your result

Validation

Builders/consumers get minimum needed access; one authoritative KPI model is identifiable; distribution/support owners are explicit.

14. Teach it back in your own words

Explain without reading

Explain why an app can be better than giving hundreds of people workspace access.

15. Protected review + changed fresh retry

Attempt protection

Attempt first. Changed retry introduces external guest/partner consumption and requires a policy check rather than assuming it works.

16. Evidence / mastery + exact next route

Mastery evidence

Builder-consumer architecture + distribution matrix + owner/freshness/security notes + explain-back. After saving evidence, continue to CLD6.

CLD6 - Workspaces and deployment concepts

1. Why it matters

Editing production directly makes rollback, testing and impact analysis difficult. Enterprise BI separates development and released environments and controls who can promote changes.

2. What you are learning

Design DEV/TEST/PROD workspace separation, deployment ownership, pre-release validation, deployment-pipeline/CI-CD concepts and rollback/impact checks without pretending every tenant has the same tooling.

3. Prerequisite / what you already know

CLD5 enterprise Service delivery and C2-04 lifecycle/change impact basics.

4. Picture it in your head

Mental model

Guarded conveyor belt: DEV is where change is built, TEST is where behavior/security are proved, and PROD is what consumers rely on. Promotion moves forward only after explicit checks; a failed control stops the conveyor. Direct production edits weaken traceability, rollback and impact analysis.

5. New words before we use them

Environment: Separated development/testing/production context for items and configuration. Deployment pipeline: Fabric capability that helps manage content deployment across pipeline stages/workspaces. CI/CD: Practices/tools for controlled integration, testing and delivery of changes.

6. Watch / visual source

VISUAL/SOURCE ROUTE: Required current references: Fabric CI/CD overview and deployment-pipeline documentation.

<https://learn.microsoft.com/en-us/fabric/cicd/cicd-overview>

After watching/opening: return to this lesson and reproduce the worked example from your own file/query/report. Save course-task evidence, not a screenshot of the source.

7. Follow with me once

Teacher demonstration - do this once with me, then close the example before your fresh task.

- 1 Developer changes measure/report in DEV; production content is not edited directly.
- 2 Run model/reconciliation/report/security checks in DEV and then controlled TEST/user acceptance as required.
- 3 Review impact: downstream reports/apps/refresh/security/config. Record release approver.
- 4 Deploy/promote items using the organization's supported method, such as Fabric deployment pipelines and/or Git-integrated CI/CD where appropriate.
- 5 After release, verify PROD refresh/security/KPI smoke checks and define rollback/redeployment path. Pipeline permissions and workspace roles are separate concerns.

8. What you should see / be able to say

Expected result

A controlled release path with evidence before and after production promotion.

9. Why it worked

Reasoning

Changes are separated from production consumers until validation/approval is complete. Professional upgrade: Fabric deployment pipelines and broader CI/CD/Git capabilities evolve. Use current docs at implementation; preserve environment separation, review, validation and rollback as durable controls.

10. Common beginner mistakes

1. DEV and PROD are tabs in one workbook only. | 2. Deployment pipeline automatically tests business correctness. | 3. Contributor can deploy everywhere without separate pipeline/workspace permissions. | 4. No post-deploy smoke or rollback plan.

11. If you get stuck - recovery ladder

Recovery

12. Now you do a small version alone

Fresh independent task

Fresh case: change Customer Segment logic used by 12 reports. Build DEV/TEST/PROD route and an impact/release checklist.

13. Validate your result

Validation

No production edit occurs before validation; approver, dependencies, post-deploy smoke and rollback are explicit.

14. Teach it back in your own words

Explain without reading

Explain deployment pipelines as release control, not a magical correctness checker.

15. Protected review + changed fresh retry

Attempt protection

Attempt first. Changed retry introduces a failed TEST validation and asks what must not reach PROD.

16. Evidence / mastery + exact next route

Mastery evidence

Environment map + release checklist + impact analysis + explain-back. After saving evidence, continue to CLD7.

CLD7 - Refresh, gateways and scheduled operations

1. Why it matters

A report can be perfect at 9 PM and wrong at 9 AM if refresh credentials, gateway mapping or schedules fail. Freshness is an operational contract, not a checkbox.

2. What you are learning

Decide when an on-premises gateway is required; distinguish cloud/on-prem connections; design credentials, refresh cadence, failure ownership and freshness checks; avoid personal-account dependency.

3. Prerequisite / what you already know

CLD2 identity, CLD5 Service and C2-04 lifecycle basics.

4. Picture it in your head

Mental model

Delivery route: cloud service -> stored credentials/identity -> gateway when the source is on-premises -> source -> refresh job -> semantic/report freshness SLA. A green job proves a job completed, not that every required upstream partition was complete or that business data is current.

5. New words before we use them

On-premises data gateway: Software bridge that lets cloud services securely reach supported data sources in a private/on-premises network. Scheduled refresh: Configured recurring refresh of a semantic model/data source connection. Freshness SLA: Maximum acceptable data age for the business use case.

6. Watch / visual source

VISUAL/SOURCE ROUTE: Required visual: Power BI Gateway getting-started video; official gateway/refresh docs control current behavior.

<https://learn.microsoft.com/en-us/power-bi/connect-data/refresh-scheduled-refresh>

After watching/opening: return to this lesson and reproduce the worked example from your own file/query/report. Save course-task evidence, not a screenshot of the source.

7. Follow with me once

Teacher demonstration - do this once with me, then close the example before your fresh task.

- 1 Classify data source location/connectivity. If Power BI/Fabric cannot reach the on-prem/private source directly, design a gateway/connection path.
- 2 Map all required source connections/credentials; one missing source mapping can make a gateway unavailable for the semantic model.
- 3 Choose refresh cadence from business need and source/load constraints-not 'as often as possible'.
- 4 Assign technical owner and business freshness owner; use non-personal/team-managed credentials where policy supports it.
- 5 After refresh, validate both technical success and business freshness/row/control totals. A successful refresh with stale source data is still a business failure.

8. What you should see / be able to say

Expected result

A refresh runbook with source connectivity, credential/gateway mapping, schedule, owners and post-refresh controls.

9. Why it worked

Reasoning

Network/credential/schedule/business freshness are tested separately. Professional upgrade: Current Power BI scheduled-refresh guidance separates gateway connection, credentials and schedule. Exact refresh limits/licensing vary; do not hard-code them into architecture without current verification.

10. Common beginner mistakes

1. Gateway required for every cloud source. | 2. Use analyst's personal password for production refresh. | 3. Refresh succeeded = source data current. | 4. No owner for expired credential/gateway outage.

11. If you get stuck - recovery ladder

Recovery

Trace Source -> Network path -> Credential -> Gateway/connection -> Semantic model -> Schedule -> Freshness validation. Save the genuine attempt before review.

12. Now you do a small version alone

Fresh independent task

Fresh case: semantic model uses cloud SaaS + on-prem SQL + SharePoint. Decide gateway/credentials and create an 8 AM freshness runbook.

13. Validate your result

Validation

Each source has correct connectivity method; owners and failure escalation exist; business freshness is checked after technical refresh.

14. Teach it back in your own words

Explain without reading

Explain why a gateway is about connectivity to a source, not a speed booster.

15. Protected review + changed fresh retry

Attempt protection

Attempt first. Changed retry simulates expired SQL credential at 06:00 and requires recovery/escalation.

16. Evidence / mastery + exact next route

Mastery evidence

Connectivity map + refresh schedule/owner matrix + failure runbook + explain-back. After saving evidence, continue to CLD8.

CLD8 - Permissions and governance

1. Why it matters

Giving people access is not one checkbox. Workspace roles, item/semantic-model access, row-level security, OneLake/item security and governance/catalog signals solve different parts of the problem.

2. What you are learning

Apply least privilege across workspace/item/data layers; use groups; distinguish access control from RLS; add sensitivity/endorsement/catalog/governance responsibilities; create periodic access review.

3. Prerequisite / what you already know

CLD2 identity, CLD5 Service and C2-04 RLS.

4. Picture it in your head

Mental model

Building access model: workspace roles grant broad collaboration privileges, item/app permissions shape consumption, and data-level controls such as RLS narrow what an authorized consumer can see. Least privilege uses groups/service principals and the smallest scope/time needed; governance adds ownership, lineage, sensitivity and auditability.

5. New words before we use them

Workspace role: Admin, Member, Contributor or Viewer permission level within a Fabric workspace. Item-level access: Permission granted to a specific report, semantic model or Fabric item rather than whole workspace. Sensitivity label: Governance/compliance classification applied under organizational policy to help protect data/content.

6. Watch / visual source

VISUAL/SOURCE ROUTE: Required current references: workspace roles, permission model, OneLake catalog Govern/Purview.

<https://learn.microsoft.com/en-us/fabric/fundamentals/roles-workspaces>

After watching/opening: return to this lesson and reproduce the worked example from your own file/query/report. Save course-task evidence, not a screenshot of the source.

7. Follow with me once

Teacher demonstration - do this once with me, then close the example before your fresh task.

- 1 Builders belong to appropriate security groups and get only required workspace roles.
- 2 Consumers receive app/item/semantic-model access rather than Contributor by convenience.
- 3 RLS limits rows inside the model for users who can query it; it does not replace workspace/item permission design.
- 4 Use OneLake catalog/endorsement/sensitivity/governance features to improve discovery, trust and policy visibility where the organization enables them.
- 5 Schedule access review: leavers, role changes, dormant broad access, external users and certified-content owners.

8. What you should see / be able to say

Expected result

Layered least-privilege design with access owner and recurring review.

9. Why it worked

Reasoning

Each security/governance control has a separate purpose and ownership. Professional upgrade: Current Fabric workspace roles are Admin/Member/Contributor/Viewer, while Fabric/Power BI also support item-specific and data-security mechanisms. Current PL-300 directly assesses workspace/item/semantic-model access, RLS and sensitivity labels.

10. Common beginner mistakes

1. RLS means user can safely be workspace Admin. | 2. Viewer and app audience are identical concepts. | 3. Certified content means anyone may access it. | 4. Permissions assigned individually forever without group/review strategy.

11. If you get stuck - recovery ladder

Recovery

Create columns: identity/group, resource, action, workspace role, item permission, data restriction, approval owner, review date. Save the genuine attempt before review.

12. Now you do a small version alone

Fresh independent task

Fresh case: Finance builders, regional managers, executives and auditors need different actions/data. Build least-privilege matrix and quarterly review checklist.

13. Validate your result

Validation

No persona gets broader access than needed; RLS and item access are not conflated; owner/review date exists.

14. Teach it back in your own words

Explain without reading

Explain workspace permission versus RLS using 'can open the building' versus 'which records can you see once inside'.

15. Protected review + changed fresh retry

Attempt protection

Attempt first. Changed retry removes a regional manager and adds an external auditor under explicit policy review.

16. Evidence / mastery + exact next route

Mastery evidence

Least-privilege matrix + governance/catalog plan + access-review checklist + explain-back. After saving evidence, continue to CLD9.

CLD9 - Monitoring and failure awareness

1. Why it matters

Enterprise analytics is not finished when the pipeline/report is published. Someone must detect failures, know the affected consumers, restore service and prevent recurrence.

2. What you are learning

Use run status/history and monitoring concepts to triage failed refresh/pipeline jobs; distinguish symptom/root cause; design incident evidence, severity, ownership, recovery and recurrence prevention.

3. Prerequisite / what you already know

CLD7 refresh operations and C2-05 logging/validation concepts helpful but not required.

4. Picture it in your head

Mental model

Smoke alarm plus dependency map: monitoring starts with the visible symptom, then traces the first broken dependency, business impact, responsible owner, recovery action and downstream validation. A successful pipeline badge is evidence about execution status, not proof that source completeness, security or freshness is correct.

5. New words before we use them

Monitoring: Observing operational health, status, history and signals from analytics jobs/items. Incident: Unplanned event that degrades correctness, freshness, access or availability. Root cause: Underlying condition whose correction prevents or reduces recurrence, not merely the visible symptom.

6. Watch / visual source

VISUAL/SOURCE ROUTE: Required current reference: Fabric monitoring hub; optional Learn Live Data Factory pipelines video.

<https://learn.microsoft.com/en-us/fabric/admin/monitoring-hub>

After watching/opening: return to this lesson and reproduce the worked example from your own file/query/report. Save course-task evidence, not a screenshot of the source.

7. Follow with me once

Teacher demonstration - do this once with me, then close the example before your fresh task.

- 1 Detect: report freshness check fails at 08:05. Record incident start, affected KPI/audience and required freshness SLA.
- 2 Check monitoring/refresh/job history: determine which job failed and capture exact status/error/time.
- 3 Trace dependencies: source availability -> credential/gateway -> ingestion/pipeline -> warehouse/model -> report. Find the first broken stage.
- 4 Recover safely: fix/renew/retry from appropriate step, then rerun downstream validation rather than only refreshing the final report.
- 5 Close incident with root cause, impact, evidence, owner and prevention action such as alert, credential rotation process, schema contract or capacity review.

8. What you should see / be able to say

Expected result

A repeatable incident runbook that connects monitoring signal to affected business output and safe recovery.

9. Why it worked

Reasoning

You diagnose dependencies from first broken layer instead of randomly clicking Refresh. Professional upgrade: Fabric's Monitor/monitoring hub currently centralizes job execution health/history. Teams may add alerts/SLAs/ticketing outside Fabric; the course assesses the operational thinking, not one UI.

10. Common beginner mistakes

1. Retry until green without root cause. | 2. Only record screenshot of final report. | 3. No business-impact/severity field. | 4. Fix pipeline but never revalidate downstream KPI/security/freshness.

11. If you get stuck - recovery ladder

Recovery

Use five boxes: Impact -> First failed job -> Dependency cause -> Recovery -> Prevention. Do not skip impact/validation. Save the genuine attempt before review.

12. Now you do a small version alone

Fresh independent task

Fresh case: nightly pipeline succeeds but semantic-model refresh fails; build incident timeline, root-cause tree and recovery/prevention actions.

13. Validate your result

Validation

Incident record includes status/error/time, affected consumers/KPI, first broken dependency, safe rerun path and prevention owner.

14. Teach it back in your own words

Explain without reading

Explain why 'Refresh failed' is a symptom, not automatically the root cause.

15. Protected review + changed fresh retry

Attempt protection

Attempt first. Changed retry uses a pipeline that succeeds technically but produces stale source data.

16. Evidence / mastery + exact next route

Mastery evidence

Incident report + dependency tree + recovery validation + prevention action + explain-back. After saving evidence, continue to CLD10.

CLD10 - Analyst - analytics engineer - data engineer collaboration

1. Why it matters

A senior analyst should not throw vague tickets at engineers or silently rebuild engineering logic in Power BI. Strong teams agree on ownership, grains, freshness, definitions and failure evidence.

2. What you are learning

Distinguish typical responsibilities without rigid job-title stereotypes; write usable data contracts/issues/change requests; identify when the analyst should self-serve, escalate or collaborate.

3. Prerequisite / what you already know

All C2-07 plus C2-06 lineage/quality and C2-05 reproducibility concepts.

4. Picture it in your head

Mental model

Contract handoff: the analyst owns the business question, KPI meaning and acceptance evidence; the analytics engineer owns reusable transformation/semantic contracts; the data engineer owns reliable source ingestion and platform pipelines. Lineage, interfaces, owners and testable acceptance criteria keep handoffs from becoming guesswork.

5. New words before we use them

Data contract: Explicit agreement about dataset structure, grain, keys, freshness, quality expectations and ownership. Handoff: Transfer of a requirement/change/incident with enough evidence for the receiving role to act. Escalation: Raising a problem to the role/team that owns the affected system or decision.

6. Watch / visual source

VISUAL/SOURCE ROUTE: Book-led; current PL-300 audience profile and Fabric collaboration context support the lesson.

<https://learn.microsoft.com/en-us/credentials/certifications/resources/study-guides/pl-300>

After watching/opening: return to this lesson and reproduce the worked example from your own file/query/report. Save course-task evidence, not a screenshot of the source.

7. Follow with me once

Teacher demonstration - do this once with me, then close the example before your fresh task.

- 1 Analyst proves impact: Revenue KPI changed after 07:00 refresh; records expected/observed values and affected report/consumers.
- 2 Trace lineage to staging table where RegionCode renamed/changed. Do not immediately hard-code a workaround in the report.
- 3 Data/analytics engineering owners receive a reproducible issue: source/object, schema diff, grain/key, first failing transformation, sample IDs, freshness/SLA and desired contract.
- 4 Engineering repairs or intentionally changes upstream contract; analyst validates business semantics and downstream KPI after redeploy.
- 5 Team records ownership/prevention: schema-change notification, tests, versioned contract or consumer impact process.

8. What you should see / be able to say

Expected result

A handoff that another technical role can act on without a meeting just to discover basic facts.

9. Why it worked

Reasoning

Business impact and technical evidence meet at a shared contract rather than being thrown over a team boundary. Professional upgrade: Current PL-300 audience guidance explicitly expects Power BI analysts to collaborate with analytics engineers and data engineers to identify/acquire data. Exact job boundaries vary by company; contract/evidence quality is the portable skill.

10. Common beginner mistakes

1. Analyst = dashboards only, engineer = everything technical. | 2. Opening 'data is wrong' ticket with no expected/observed evidence. | 3. Rebuilding official transformations locally to avoid collaboration. | 4. Engineer owns KPI business definition by default.

11. If you get stuck - recovery ladder

Recovery

Use issue template: Business impact -> expected/observed -> dataset/item -> grain/key -> lineage point -> freshness -> evidence -> requested action -> owner/deadline. Save the genuine attempt before review.

12. Now you do a small version alone

Fresh independent task

Fresh case: Support SLA dashboard needs a new event field unavailable downstream. Decide which role owns definition, ingestion, transformation, semantic exposure and business validation.

13. Validate your result

Validation

Every deliverable/decision has one accountable owner; handoff includes reproducible evidence and agreed acceptance check.

14. Teach it back in your own words

Explain without reading

Explain analyst, analytics engineer and data engineer collaboration as overlapping responsibilities connected by contracts-not three isolated ladders.

15. Protected review + changed fresh retry

Attempt protection

Attempt first. Changed retry is a schema-deprecation notice requiring impact analysis and coordinated migration.

16. Evidence / mastery + exact next route

Mastery evidence

RACI/ownership map + issue/change template + acceptance validation + explain-back. After saving evidence, complete the module mastery checklist and follow the Gate A / completion route.

C2-07 mastery checklist + quick reference

Use after the first pass; live UI exposure never replaces reasoning

TEN QUESTIONS

CLD1 Can you separate tenant/workspace/capacity/items? | CLD2 Is failure compute, identity, credential or data permission? | CLD3 Why does each Fabric workload exist? | CLD4 Who owns authoritative OneLake items and why copy/shortcut? | CLD5 How do builders and consumers differ? | CLD6 What prevents an unvalidated change reaching PROD? | CLD7 Does every source have connectivity/credential/refresh/freshness ownership? | CLD8 Which workspace/item/data/governance control grants or restricts access? | CLD9 What is the first broken dependency and recovery/prevention? | CLD10 Can another role act on your handoff without rediscovering the problem?

BLOCK MASTERY IF

Admin everywhere; tenant/workspace/capacity confused; every Fabric workload added; uncontrolled data copies; consumers given build roles; direct PROD edits; personal refresh credentials; gateway misunderstood; RLS treated as all security; incident only retried; vague 'data wrong' engineering ticket.

FAST OPERATIONAL DEBUG

Identity -> resource/action -> workspace/item/data permission -> source connectivity/credential -> job/refresh -> monitoring history -> business freshness/output -> owner/incident/prevention. Find first broken layer.

OPTIONAL PATH

C2-07 is optional Senior/BI enterprise depth. It overlaps heavily with later engineering operations, but direct Course 1 -> C2B -> C3 learners receive required cloud/engineering prerequisites in their own route.

C2-07 source / video registry

Current Microsoft docs control changing platform behavior

CLD1-4 - What is Microsoft Fabric?

<https://learn.microsoft.com/en-us/fabric/get-started/microsoft-fabric-overview>

Current SaaS/workloads/OneLake overview; updated 2026-03-31

CLD4 - What is OneLake?

<https://learn.microsoft.com/en-us/fabric/onelake/onelake-overview>

Tenant-wide logical lake; current 2026 reference

CLD3-4 - Microsoft OneLake documentation

<https://learn.microsoft.com/en-us/fabric/onelake/>

OneLake security/catalog/shortcuts navigation

CLD5 - Collaborate and share Power BI content

<https://learn.microsoft.com/en-us/power-bi/collaborate-share/collaborate-share-overview>

Workspaces/apps/share/endorsement

CLD5-6 - Workspaces in Power BI

<https://learn.microsoft.com/en-us/power-bi/collaborate-share/service-new-workspaces>

Current enterprise workspace reference

CLD6,8 - Roles in Fabric workspaces

<https://learn.microsoft.com/en-us/fabric/fundamentals/roles-workspaces>

Admin / Member / Contributor / Viewer

CLD6 - Fabric CI/CD overview

<https://learn.microsoft.com/en-us/fabric/cicd/cicd-overview>

Current release/source-control concepts

CLD6 - Deployment pipelines

<https://learn.microsoft.com/en-us/fabric/cicd/deployment-pipelines/intro-to-deployment-pipelines>

Managed stage deployment concepts

CLD7 - Configure scheduled refresh

<https://learn.microsoft.com/en-us/power-bi/connect-data/refresh-scheduled-refresh>

Gateway connection, credentials and schedules

CLD7 - On-premises data gateway

<https://learn.microsoft.com/en-us/data-integration/gateway/service-gateway-onprem>

Current gateway concept

CLD7 - Gateway visual

<https://www.youtube.com/watch?v=Vlv1tbuZNcM>

Optional visual; official docs control current behavior

CLD8 - Fabric permission model

<https://learn.microsoft.com/en-us/fabric/security/permission-model>

Fabric access/security model

CLD8 - Govern Fabric data with OneLake catalog

<https://learn.microsoft.com/en-us/fabric/governance/onelake-catalog-govern>

Governance posture/recommended actions

CLD8 - Microsoft Purview + Fabric

<https://learn.microsoft.com/en-us/fabric/governance/microsoft-purview-fabric>

Governance/security/compliance integration

CLD9 - Fabric monitoring hub

<https://learn.microsoft.com/en-us/fabric/admin/monitoring-hub>

Centralized job health/history

CLD5-8,10 - PL-300 study guide

<https://learn.microsoft.com/en-us/credentials/certifications/resources/study-guides/pl-300>

Skills measured as of Apr 20 2026

CLD5-6 - PL-300 Manage and Secure video

<https://learn.microsoft.com/en-us/shows/exam-readiness-zone/preparing-for-pl-300-manage-and-secure-power-bi>

Optional exam-alignment visual

CLD3-5 - Power BI Meets Microsoft Fabric

<https://learn.microsoft.com/en-us/shows/fabric-tech-talk-fridays/power-bi-meets-microsoft-fabric-whats-in-it-for-you>

Optional platform visual

VERIFY / FALLBACK

Rechecked 2026-09-10. External media is support only; internal lessons and synthetic fixtures remain sufficient if tenant/video access fails.