

C2_09 Final Gate B - Protected Review RC4

Healthcare Network Operations - open only after the complete saved genuine Gate attempt.

Protection

Do not use this as a pre-solve answer sheet. Compare it only after your own evidence is saved. Then close the review and complete the assigned changed fresh retry.

Page 1 - Objective controls and first broken dependency

- Appointments: $120 = 60 \text{ Completed} + 30 \text{ Cancelled} + 30 \text{ Scheduled}$. A completed-status count is not automatically the same as clinically valid care or claims realization.
- Initial claims = 60 rows. Rerun = 14 rows containing 8 existing claim IDs and 6 genuinely new claims. Naive append produces 74 rows; duplicate-safe merge produces 66 unique claims.
- Naive append claim amount = 92,576.40 versus 81,158.06 after claim_id deduplication: 11,418.34 of duplicate inflation.
- After deduplication, 64 claims link to Completed appointments, 1 to Cancelled and 1 to Scheduled. Finance claims therefore cannot be used as a proxy for valid completed care without a rule.
- patient_id_sensitive is explicitly sensitive and must not appear in executive/public evidence. Clinic targets are Month-Clinic grain (16 rows).

Interpretation boundary

Exact controls test grain, completeness and rerun safety. They do not create a single automatic business recommendation; definitions, period, prioritization metric and uncertainty must be explicit.

Page 2 - Reasoning audit

Decision / definitions

Keep scheduled volume, operational completed status, clinically valid completion and submitted/paid claim realization separate. Resolve ownership instead of picking the convenient count.

SQL / grain

Use claim_id as the rerun idempotency key; test duplicate IDs before append. Appointment and claim facts stay separate; link for reconciliation but do not merge their measures into one fact grain.

Model / DAX

Clinic/date dimensions can filter appointment and claim facts independently. Measures should label appointment status and claim status explicitly.

Failure / rerun

First broken dependency is append semantics. Replace/merge/upsert the rerun by claim_id, reconcile 66 unique claims and 81,158.06 amount, then republish only after duplicate and status checks pass.

Executive / privacy

Remove patient identifiers from executive evidence. A clinic can have claims despite no valid completed-care count, so intervention logic must state which process failure is being prioritized.

Page 3 - Critical floors, repair and changed retry

- Duplicate claims after rerun
- Patient-level/sensitive identifier leakage
- Claims presented as completed clinical care without a governing rule
- Fact fan-out or mismatched Month-Clinic target grain
- No idempotent claim_id merge/reconciliation
- Causal overclaim, unsafe rerun or protected-answer contamination

Repair rule	Repair only the weak competency with a targeted exercise. Do not reopen this review while solving the changed case.
Changed fresh retry	Gate B exists only after targeted repair from Gate A. After review, close it and solve the assigned changed retry. If another critical weakness remains, Gate C is the next protected fresh assessment only when explicitly assigned.
Evidence to save	First attempt, review diagnosis, targeted repair artifact, changed-retry submission, reconciliation/validation evidence, one failure/recovery note, privacy/security check and own-words defense.